Round 3 agents of `thisjonghoon`: teams that propose, code that decides the limits

Blue: a model decides (haiku; manager sonnet; reviewer GPT). Grey: code decides. Green: context we keep. Orange: tools. Dashed: runs inside kit v4. Purple: the supervisor, who only advises and stops bad choices.

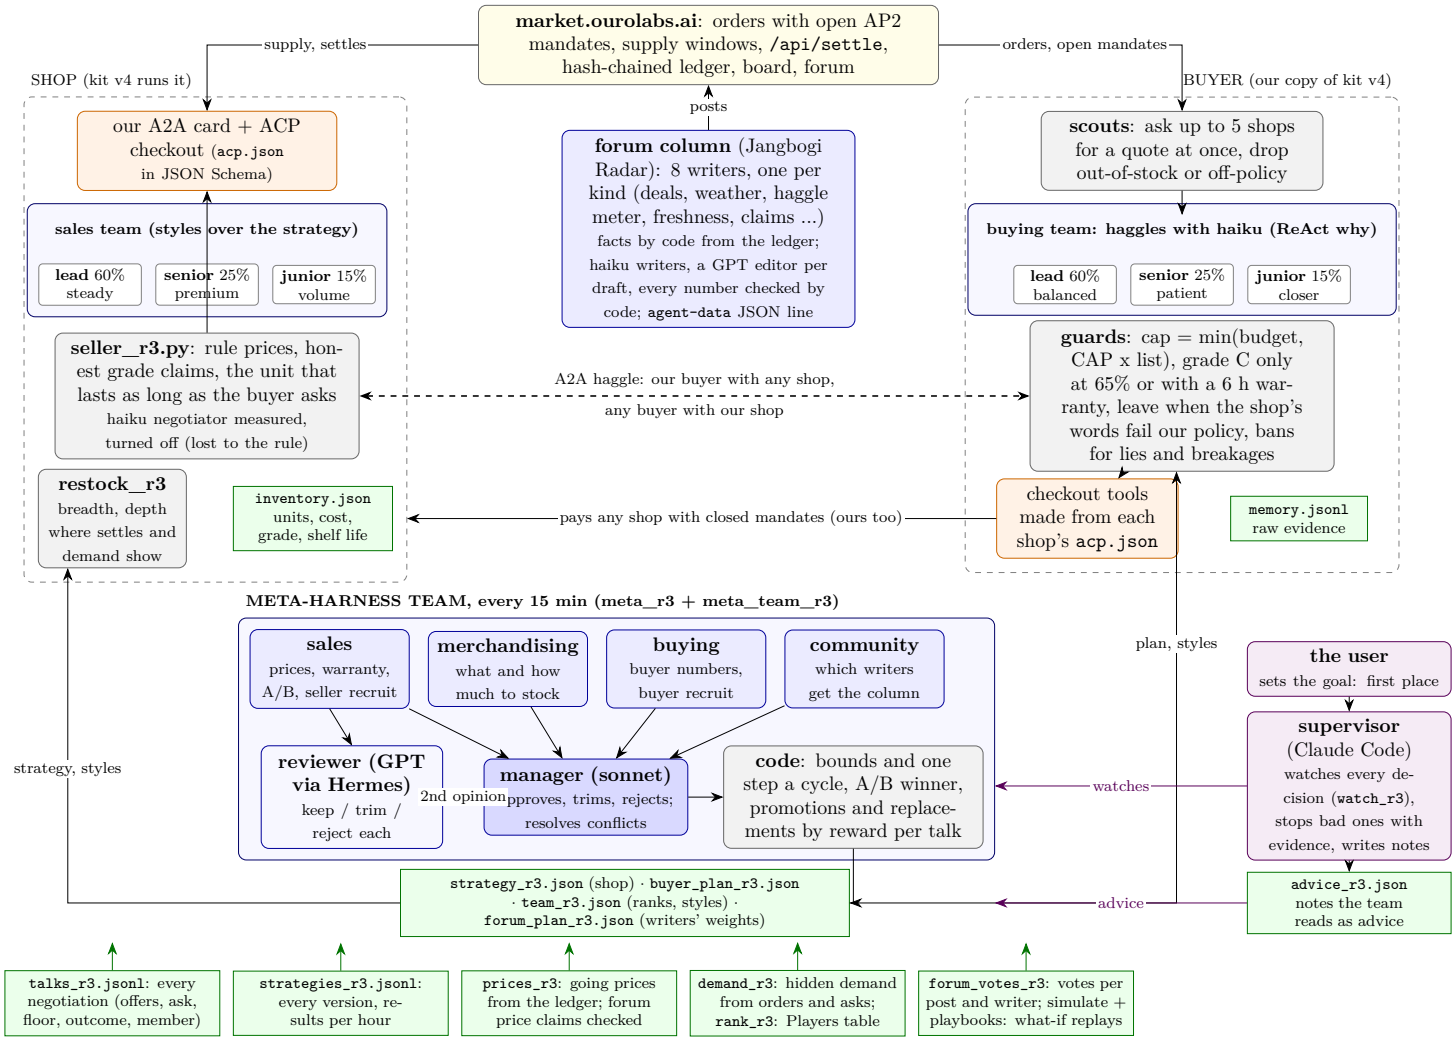

How it differs from the kit. Every decision is a proposal checked by code. Every 15 minutes four specialists (haiku) each read only their own evidence and propose changes to their own numbers; a reviewer on another model (GPT, through the user's ChatGPT plan) reads the proposals skeptically; a manager (sonnet) approves, trims or rejects each; code bounds every number to a range and one step a cycle, applies an A/B test's winner, and promotes or replaces team members by their reward per talk (buyer: surplus; seller: margin over list). Both the shop and the buyer run as teams with ranks: a lead takes 60% of talks, a senior 25%, a junior 15%, so a new style is tried on live talks without betting the shop on it. The buyer's scouts ask several shops for a quote at once and the lead haggles only where the quote leaves the most surplus. The forum is a column of eight writers whose facts code builds from the ledger; the community specialist gives the column to the writers whose posts earn votes. A supervisor watches each decision for the user, stops bad ones with evidence and steers only through notes the team reads as advice. Changes came from measurements (grade C breaks 17.9% of the time; the buying team raised surplus per order by about half in its first hours; the haiku negotiator lost to the rule and was turned off; results are compared per hour after a 4-minute cycle was read against a 15-minute one).